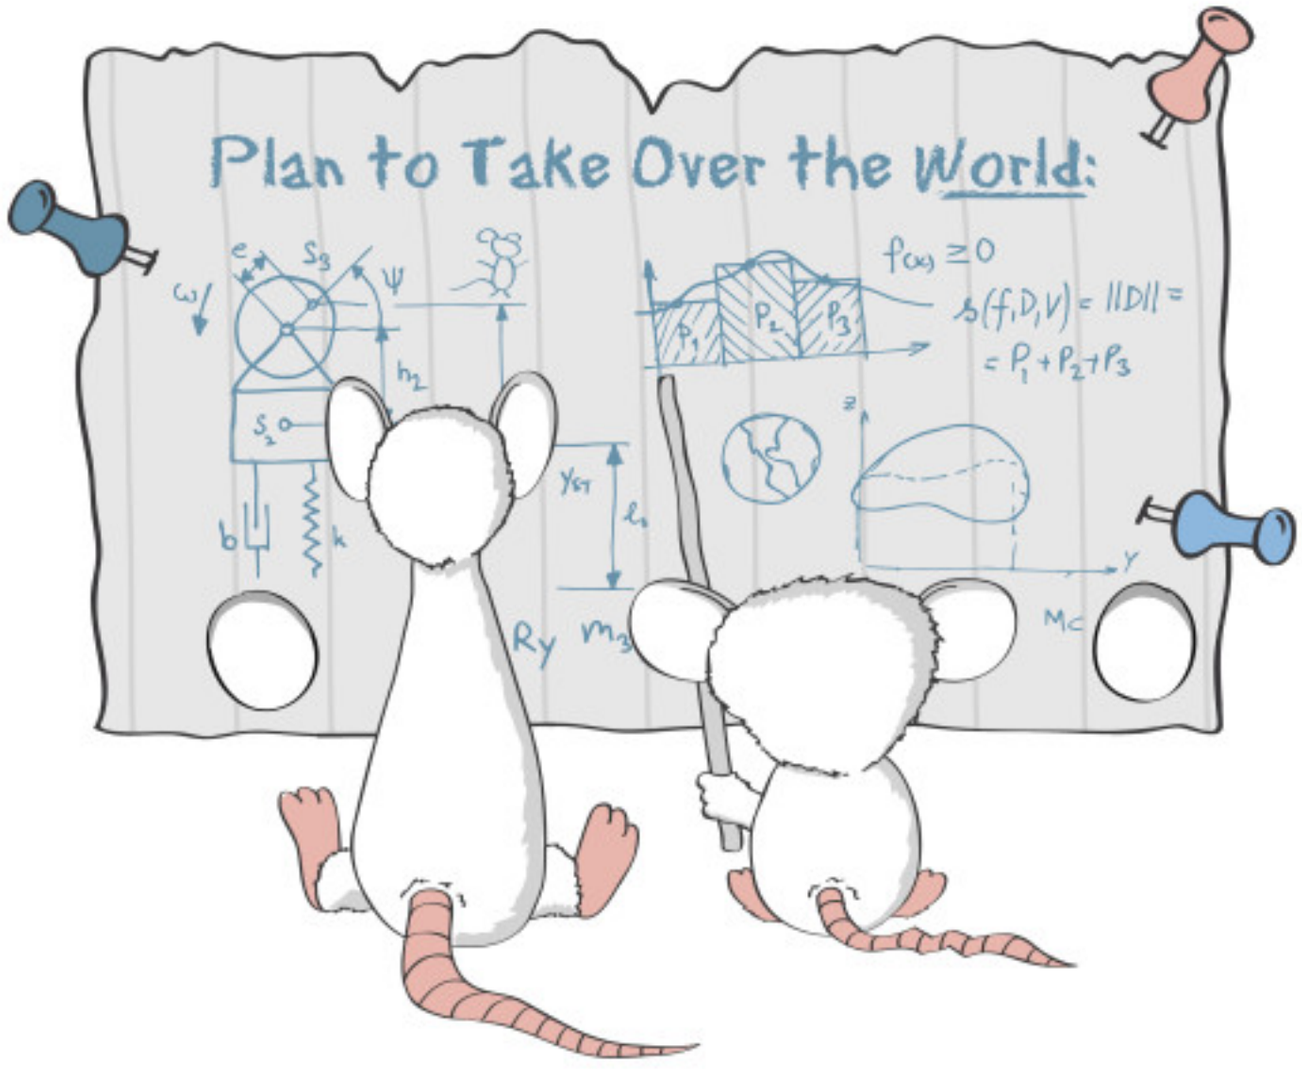

Work to Machine!

DIAS ML course by Martin Topinka

https://github.com/toastmaker/ml-dias

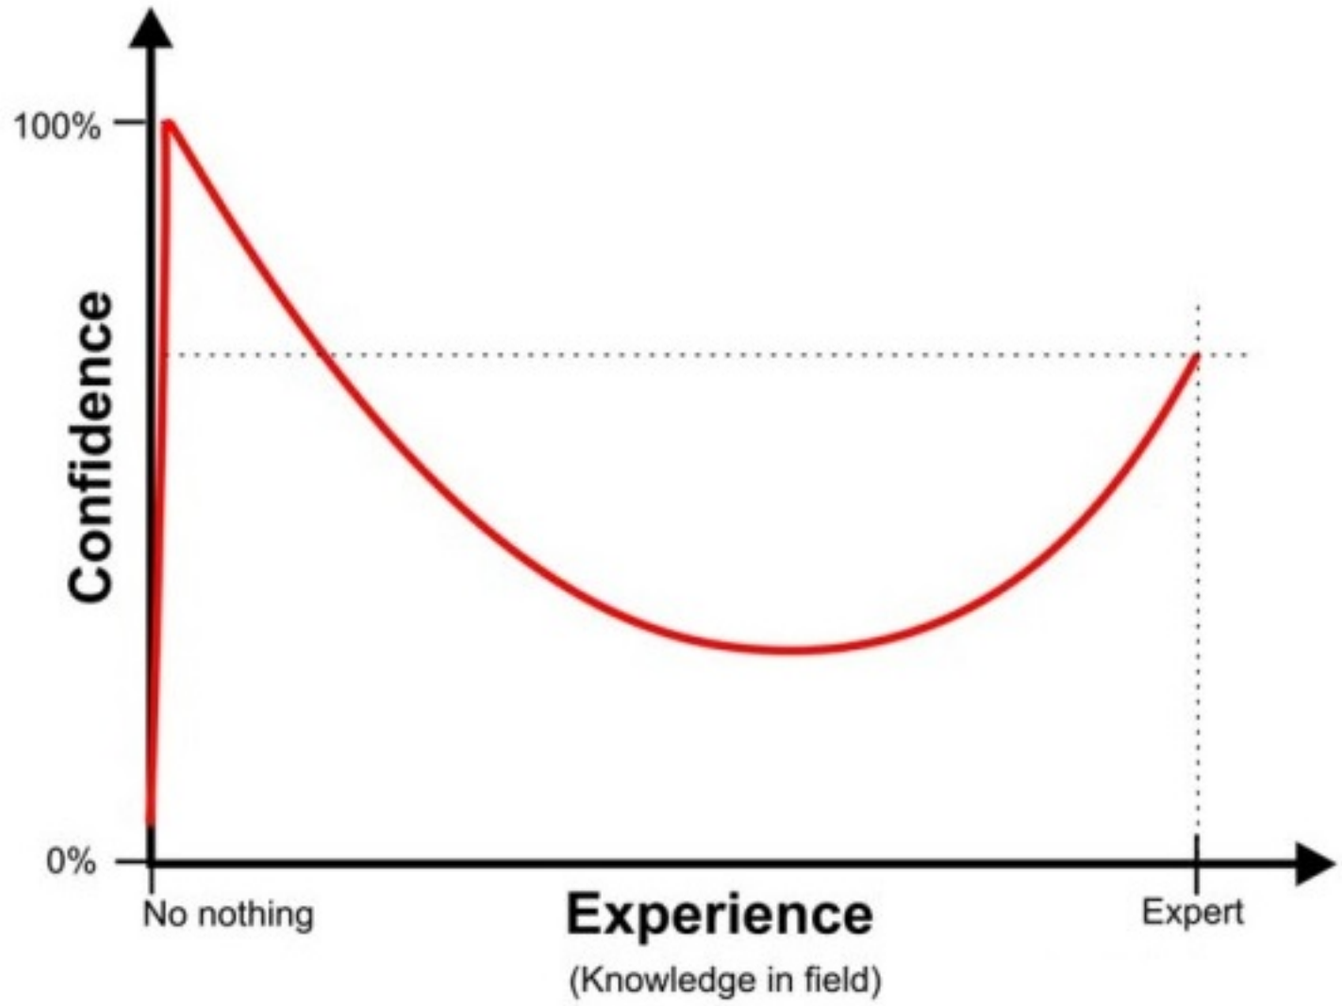
100%
Confidence
0%
No nothing
Experience
(Knowledge in field)
Expert

Introduce yourself and your goals

3Vs':

- **Volume (large archives)**
- **Velocity (continuous flow)**
- **Variety (complexity)**

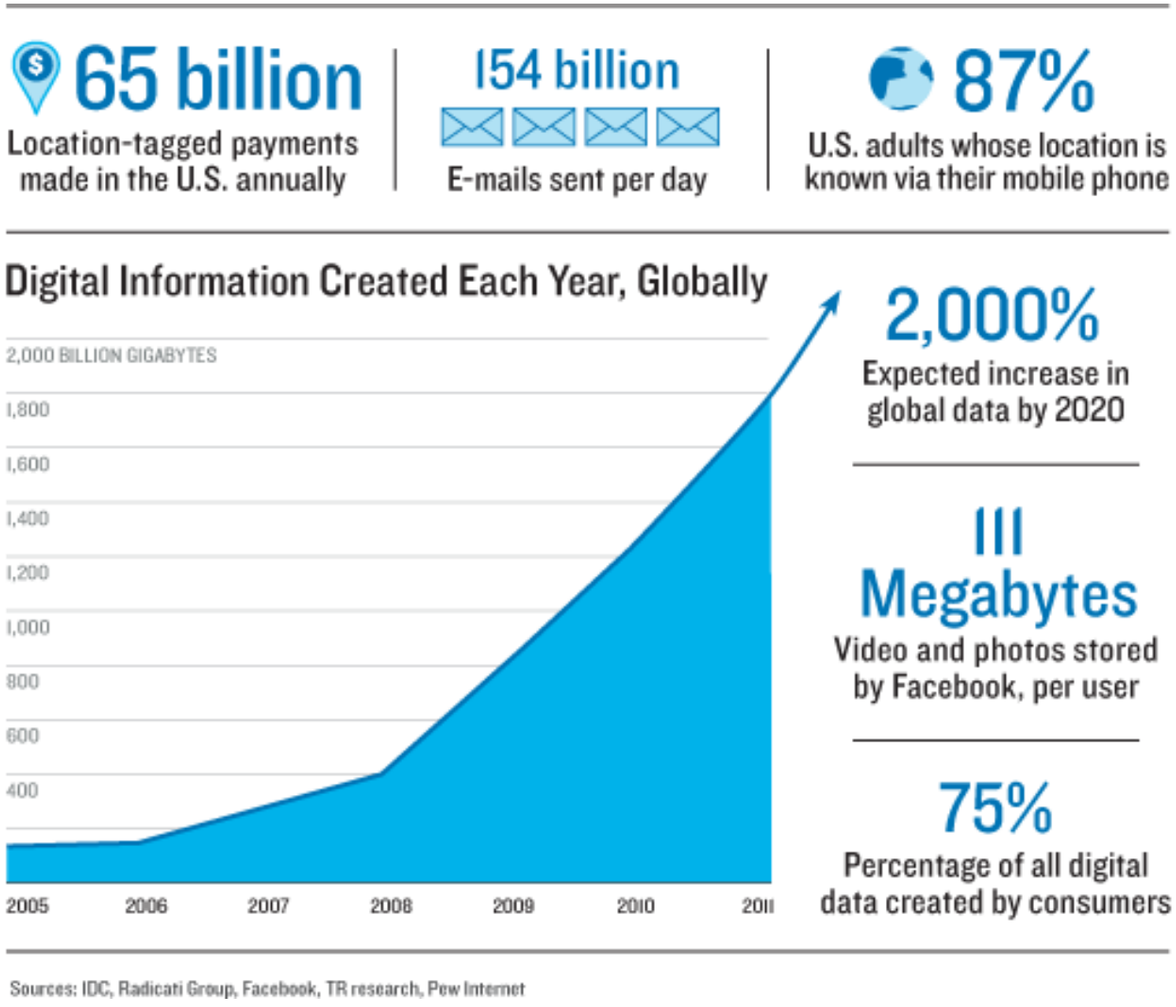

LSST

DR11 37 10^9 objects, 7 10^{12} sources,
5.5 million 3.2 Gigapixel images

30 terabytes of data nightly

Final volume of raw image data = 60 PB

Final image collection (DR11) = 0.5 EB

Final catalog size (DR11) = 15 PB

Machine…

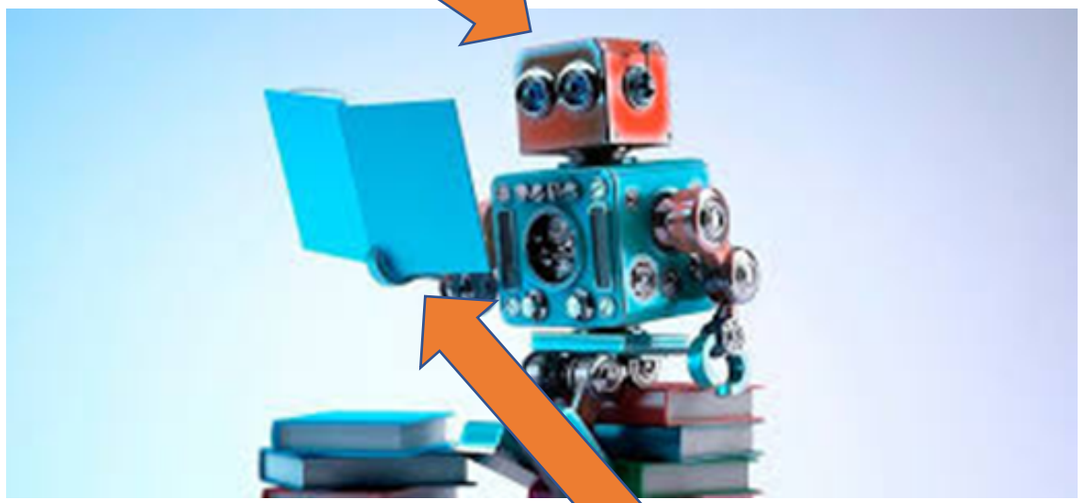

… Learning?!?

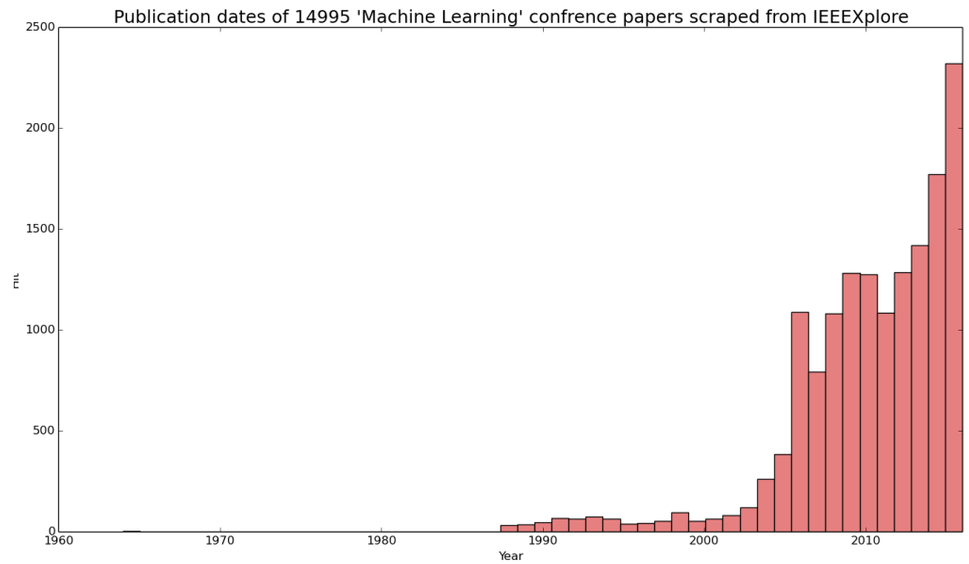

Machine learning algorithms can figure out how to perform tasks by generalising from examples (experience).

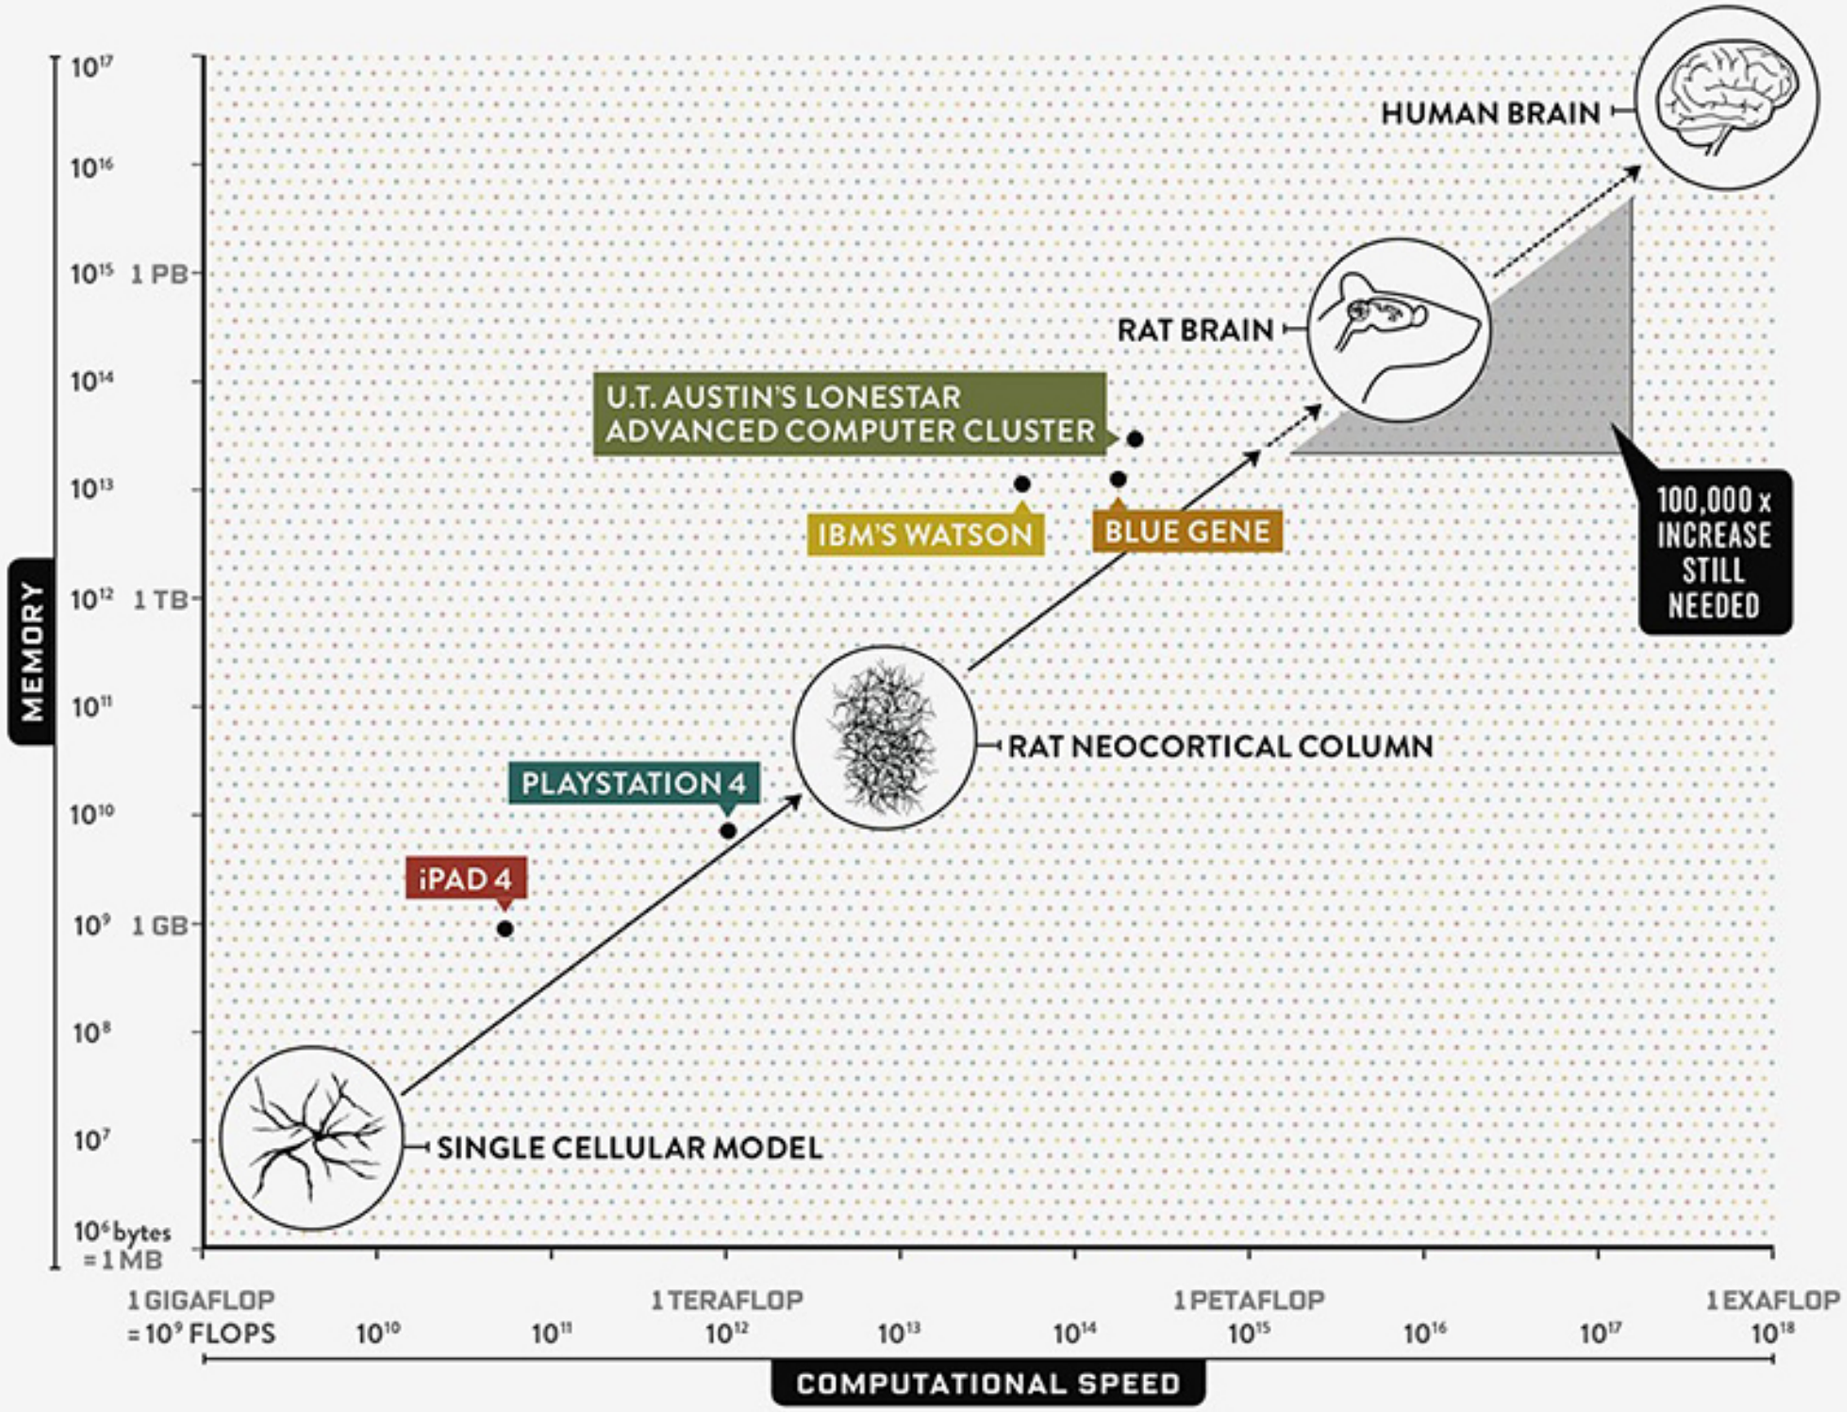
MEMORY
10^{17}
10^{16}
10^{15} 1 PB
10^{14}
10^{13}
10^{12} 1 TB
10^{11}
10^{10}
10^{9} 1 GB
10^{8}
10^{7}
10^{6} bytes = 1 MB
HUMAN BRAIN
RAT BRAIN
U.T. AUSTIN'S LONESTAR ADVANCED COMPUTER CLUSTER
IBM'S WATSON
BLUE GENE
100,000 x INCREASE STILL NEEDED
RAT NEOCORTICAL COLUMN
PLAYSTATION 4
iPAD 4
SINGLE CELLULAR MODEL
1 GIGAFLOP = 10^{9} FLOPS
10^{10}
10^{11}
1 TERAFLOP 10^{12}
10^{13}
10^{14}
1 PETAFLOP 10^{15}
10^{16}
10^{17}
1 EXAFLOP 10^{18}
COMPUTATIONAL SPEED

Name	# of neurons / # of synapses	Visuals
Caenorhabditis elegans	302	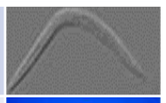
Hydra vulgaris	5,600	
Homarus americanus	100,000	
Blatta Orientalis	1,000,000	
Nile Crocodile	80,500,000	
Digital Reasoning NN (2015)	**~86,000,000 (est.) / 1.6E11**	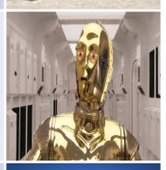
Rattus Rattatouillensis	200,000,000	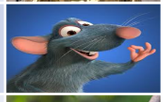
Blue and yellow macaw	1,900,000,000	
Chimpanzee	28,000,000,000	
Homo Sapiens Sapiens	**86,000,000,000 / 1.5E14**	
African Elephant	257,000,000,000	

Machine Learning in the Sky

- Machine Learning owes a lot to astronomy: least-square regression for orbital parameter estimation (Legendre-Laplace-Gauss)

Data Big Bang in Astronomy too:

10^9 object photometric catalogs from USNO, 2MASS, SDSS…

$10^{6\text{-}8}$ spectroscopic catalogs from SDSS, LAMOST…

$10^{6\text{-}7}$ multi-wavelength source catalogs from WISE, eROSITA…

10^9 object x 10^2 epochs surveys like **LSST**, DES, PTF, CRTS, SNF, VVV, Pan-STARRS, Stripe 82

Spectral-image datacubes from VLA, ALMA, IUFs, … JWST …

- **Supervised** - classification, regression (fitting)

- **Unsupervised** - clustering, graphs/trees, transformations, typically in multi-dim

- **Semi-supervised**, genetic algorithms, GANs...

Essentials of Learning

- Pattern must exist
- Mapping “target function” is unknown or expensive to calculate
- We have the data (and computing resources…)

Target function...

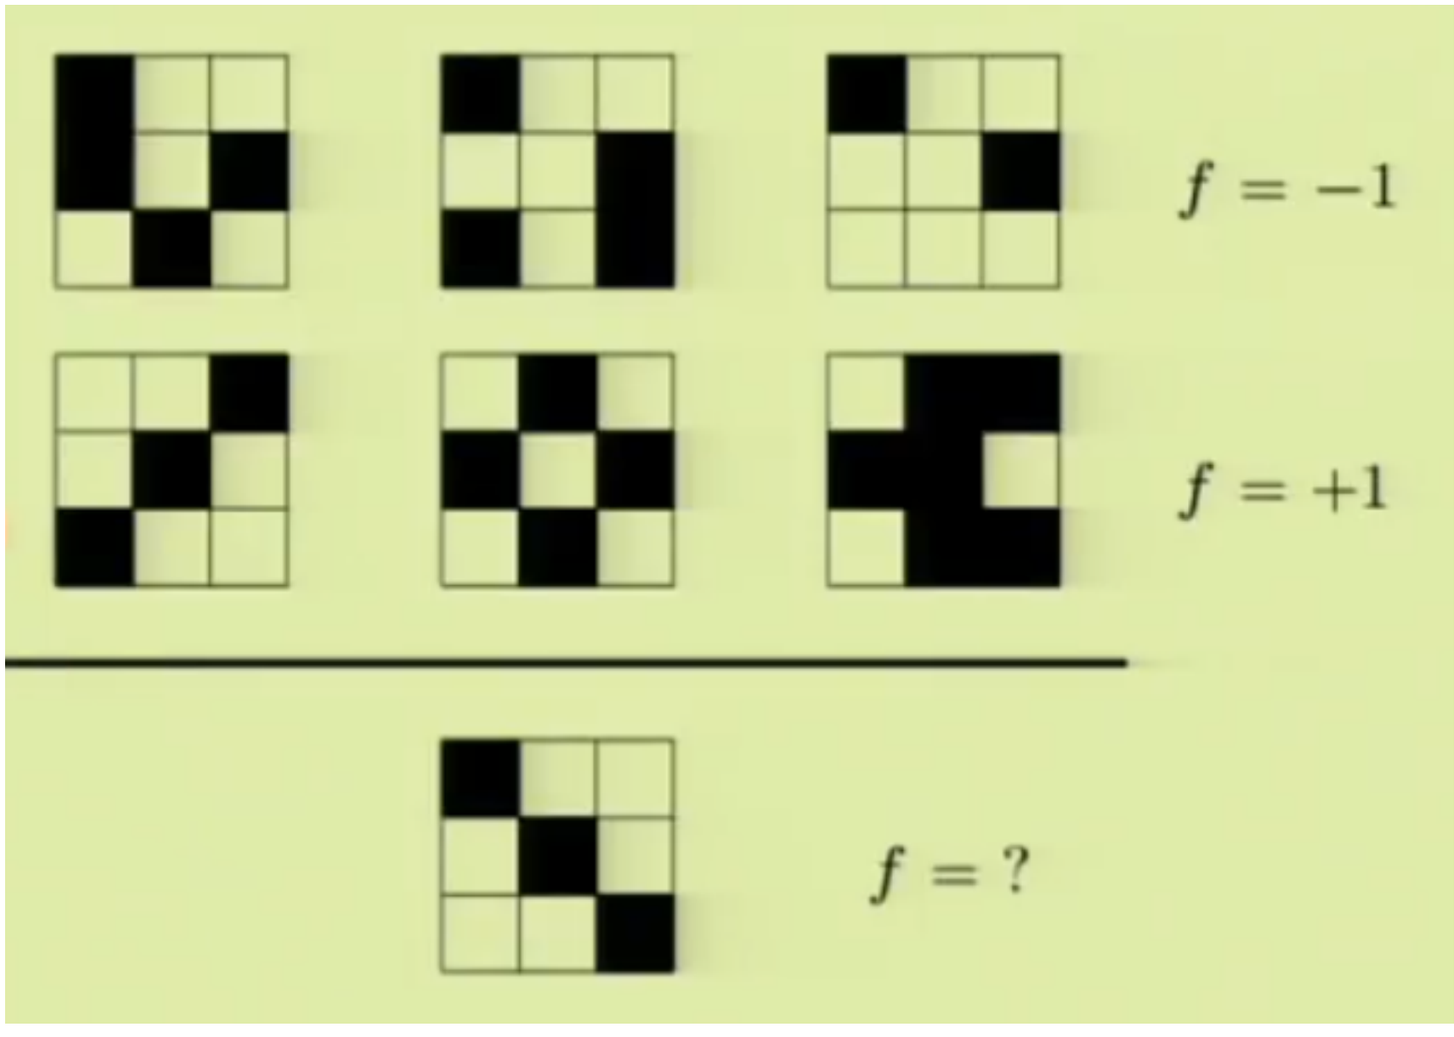

Learning Diagram

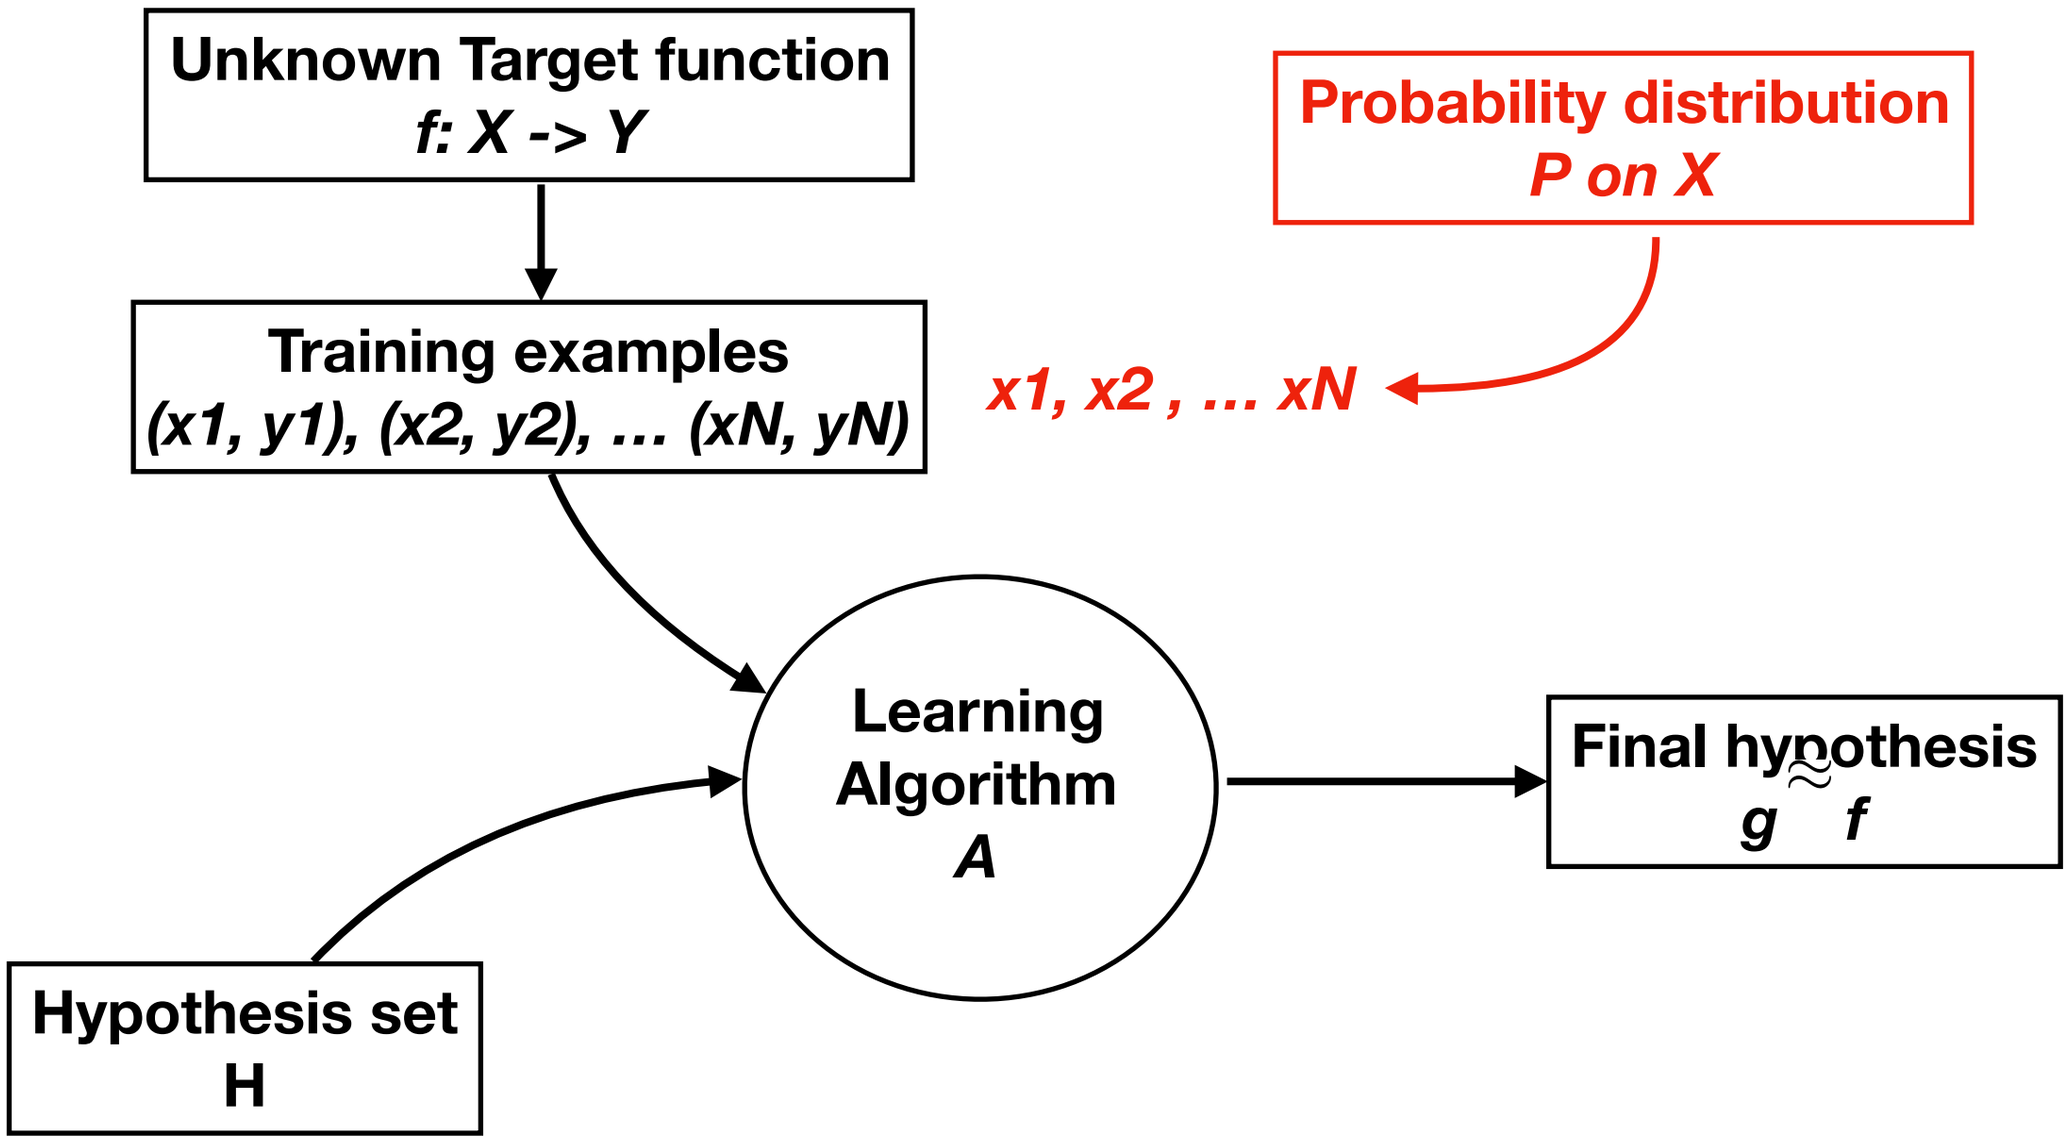

Loss function

- imbalance classes

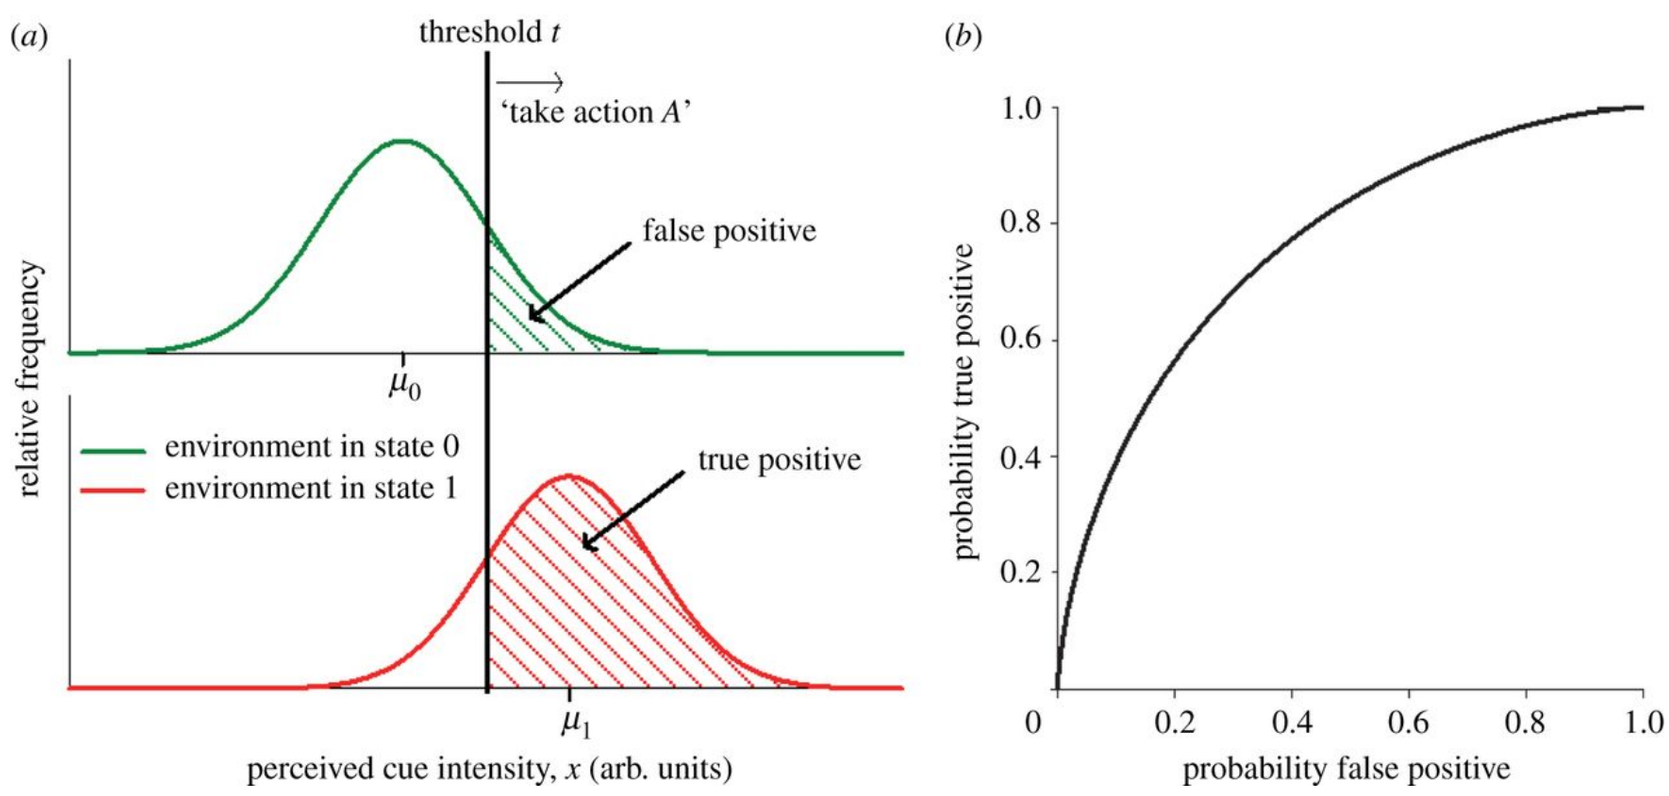

Bias - Variance Trade-off

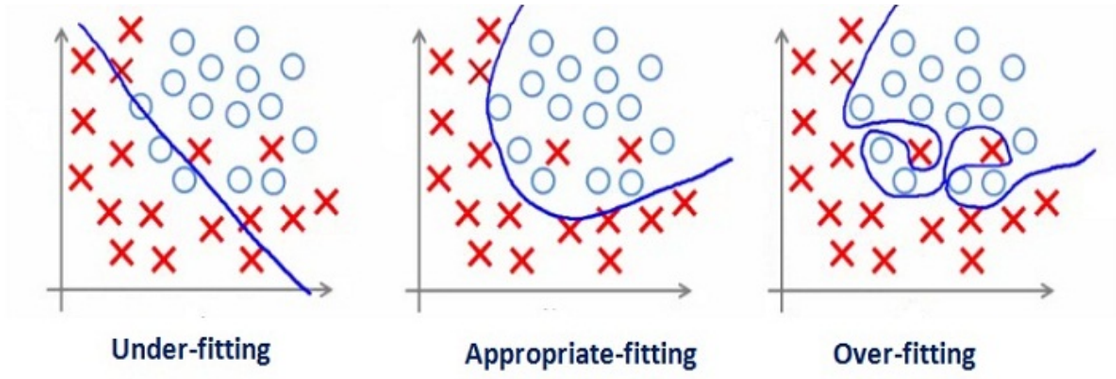

- Under-fitting/over-fitting
- in sample vs out of sample error
- VC dimension

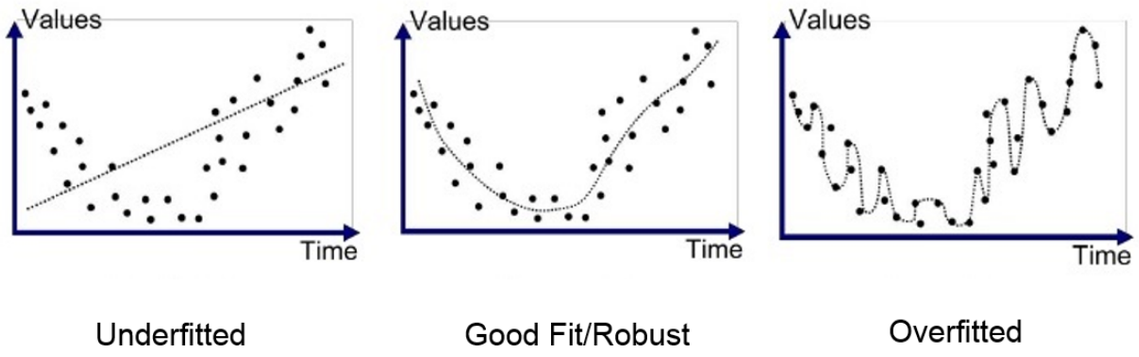

(cross)-Validation

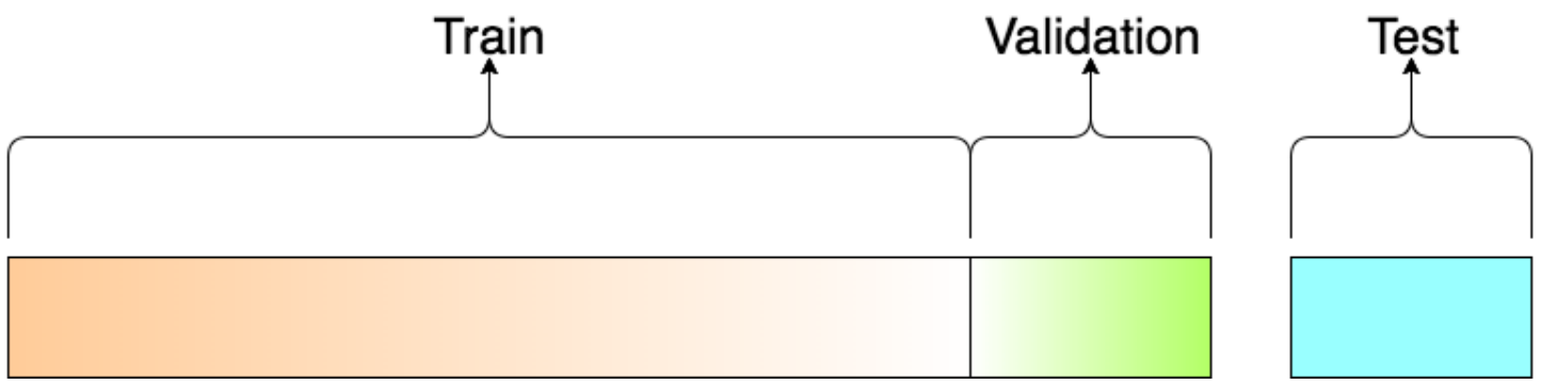

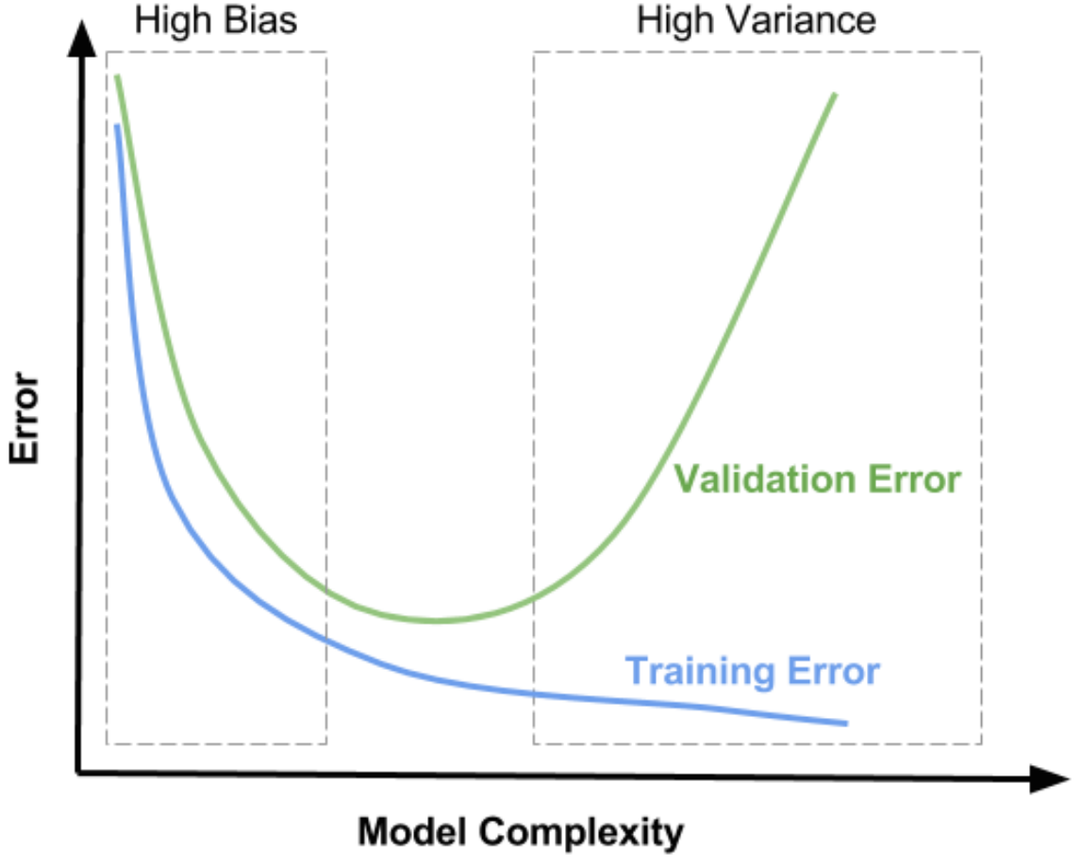

	Remedies

High Bias	• Train longer
	• Increase model complexity
	  • more features
	  • more parameters,
	  • richer architecture

High Variance	• Get more data
	• Decrease model complexity
	  • less features
	  • less parameters,
	  • simpler architecture
	• Regularization
	• Early stopping
	• Drop-out

Data Augmentation:

- When more data are needed, make up new ones! (The way of the god.)
- Translate, rotate, flip, crop, lighten/darken, add noise, dephase, etc.

Cats

Dogs

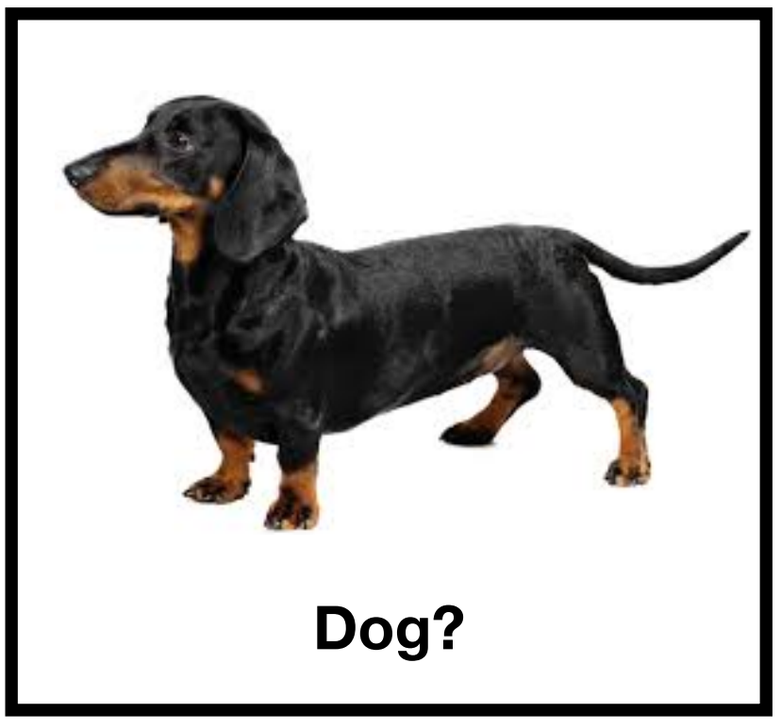

Dog?

50% dog, 50% cat?

Machines are lazy and love shortcuts

Correlations != Reasoning

Justin King
@Justinkingnews

Headline: "Women who own horses live longer"

Implied correlation: horses make you live longer.

Reality: if you own a horse, you can probably afford health insurance.

9/10/18, 10:18 PM

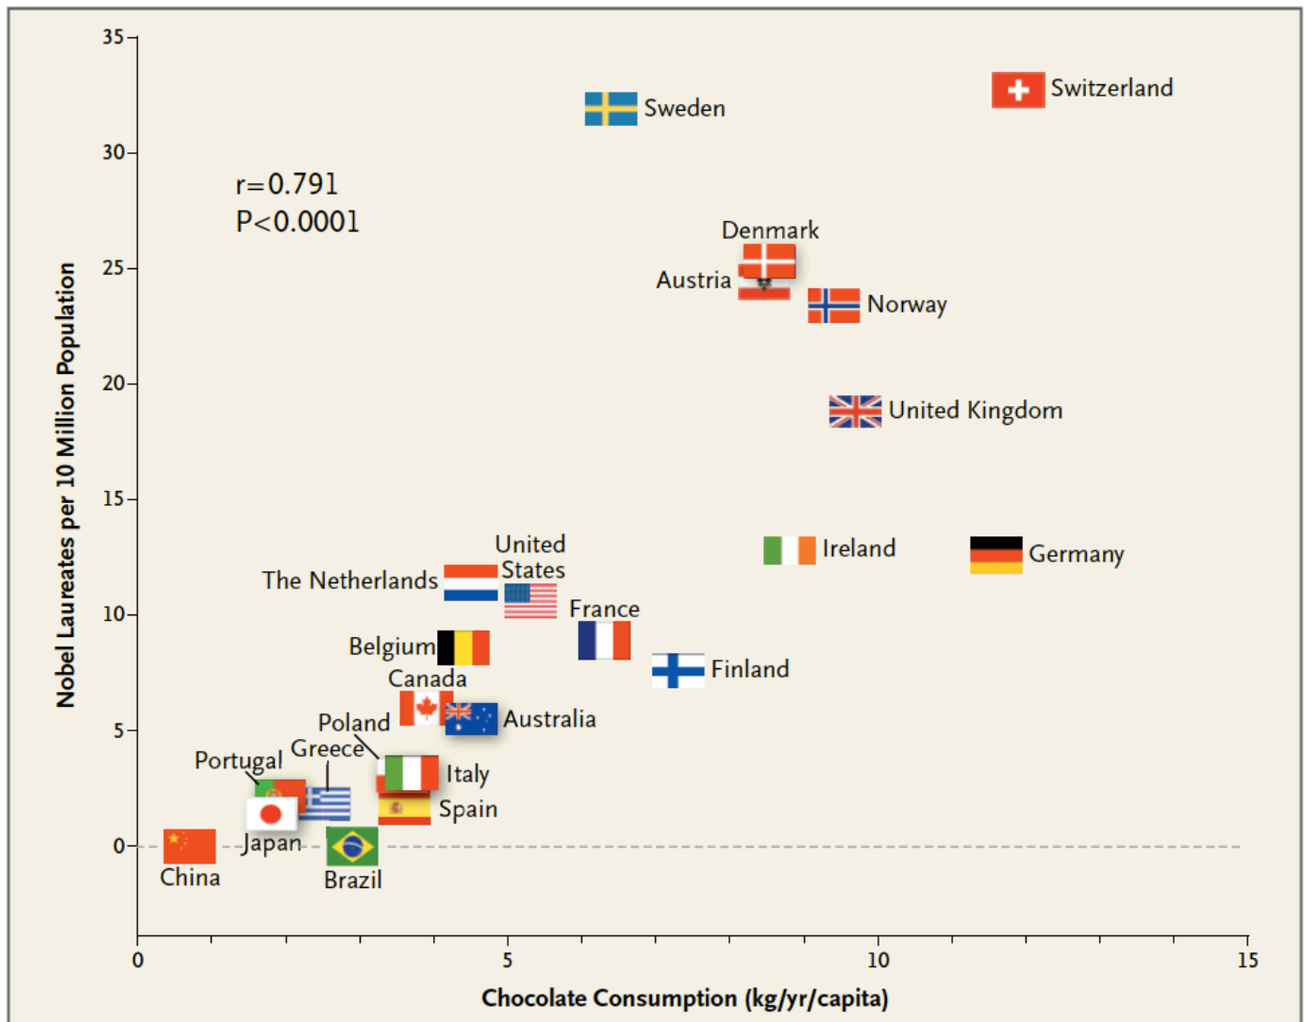

Figure 1. Correlation between Countries' Annual Per Capita Chocolate Consumption and the Number of Nobel Laureates per 10 Million Population.

Curse of Dimensionality

- Too many features
  - Expensive to store
  - Slowing down computation
  - Subject to *Dimensionality curse*
- Sample space gets harder and harder to fill as dimensions grow
- A reason why too many features lead to overfitting as data become **sparse**

- *“If people can see in multi-dimensions we would not need machine learning”*
- More and more data needed to fill the same % of space
- Distance measure degenerates, most of the points are at the surface of a sphere (distance r)

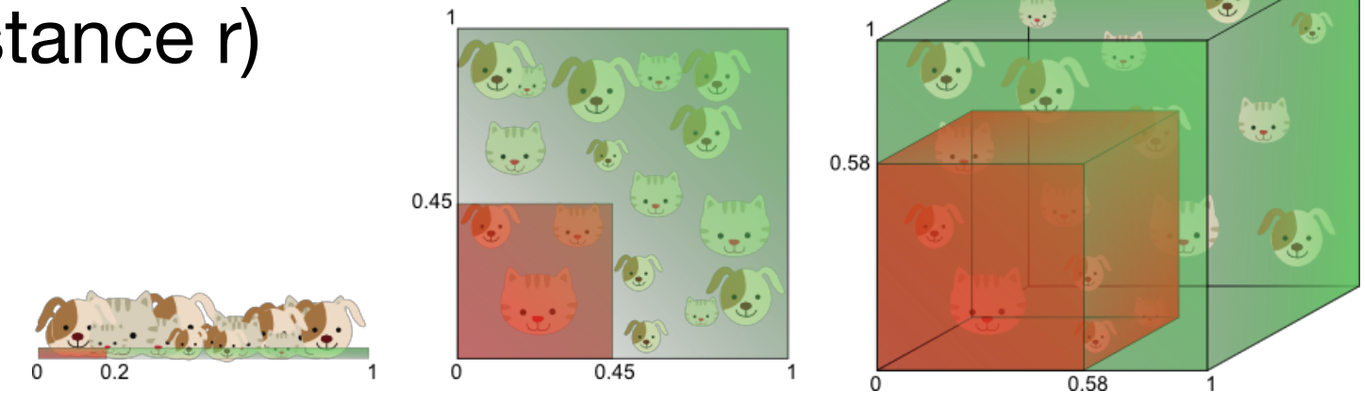

How Many Shades of Gray Can you Distinguish?

Answer…

Value encodes continuous variables (less well)

How Many Colors?

Hue encodes nominal variables

Answer…

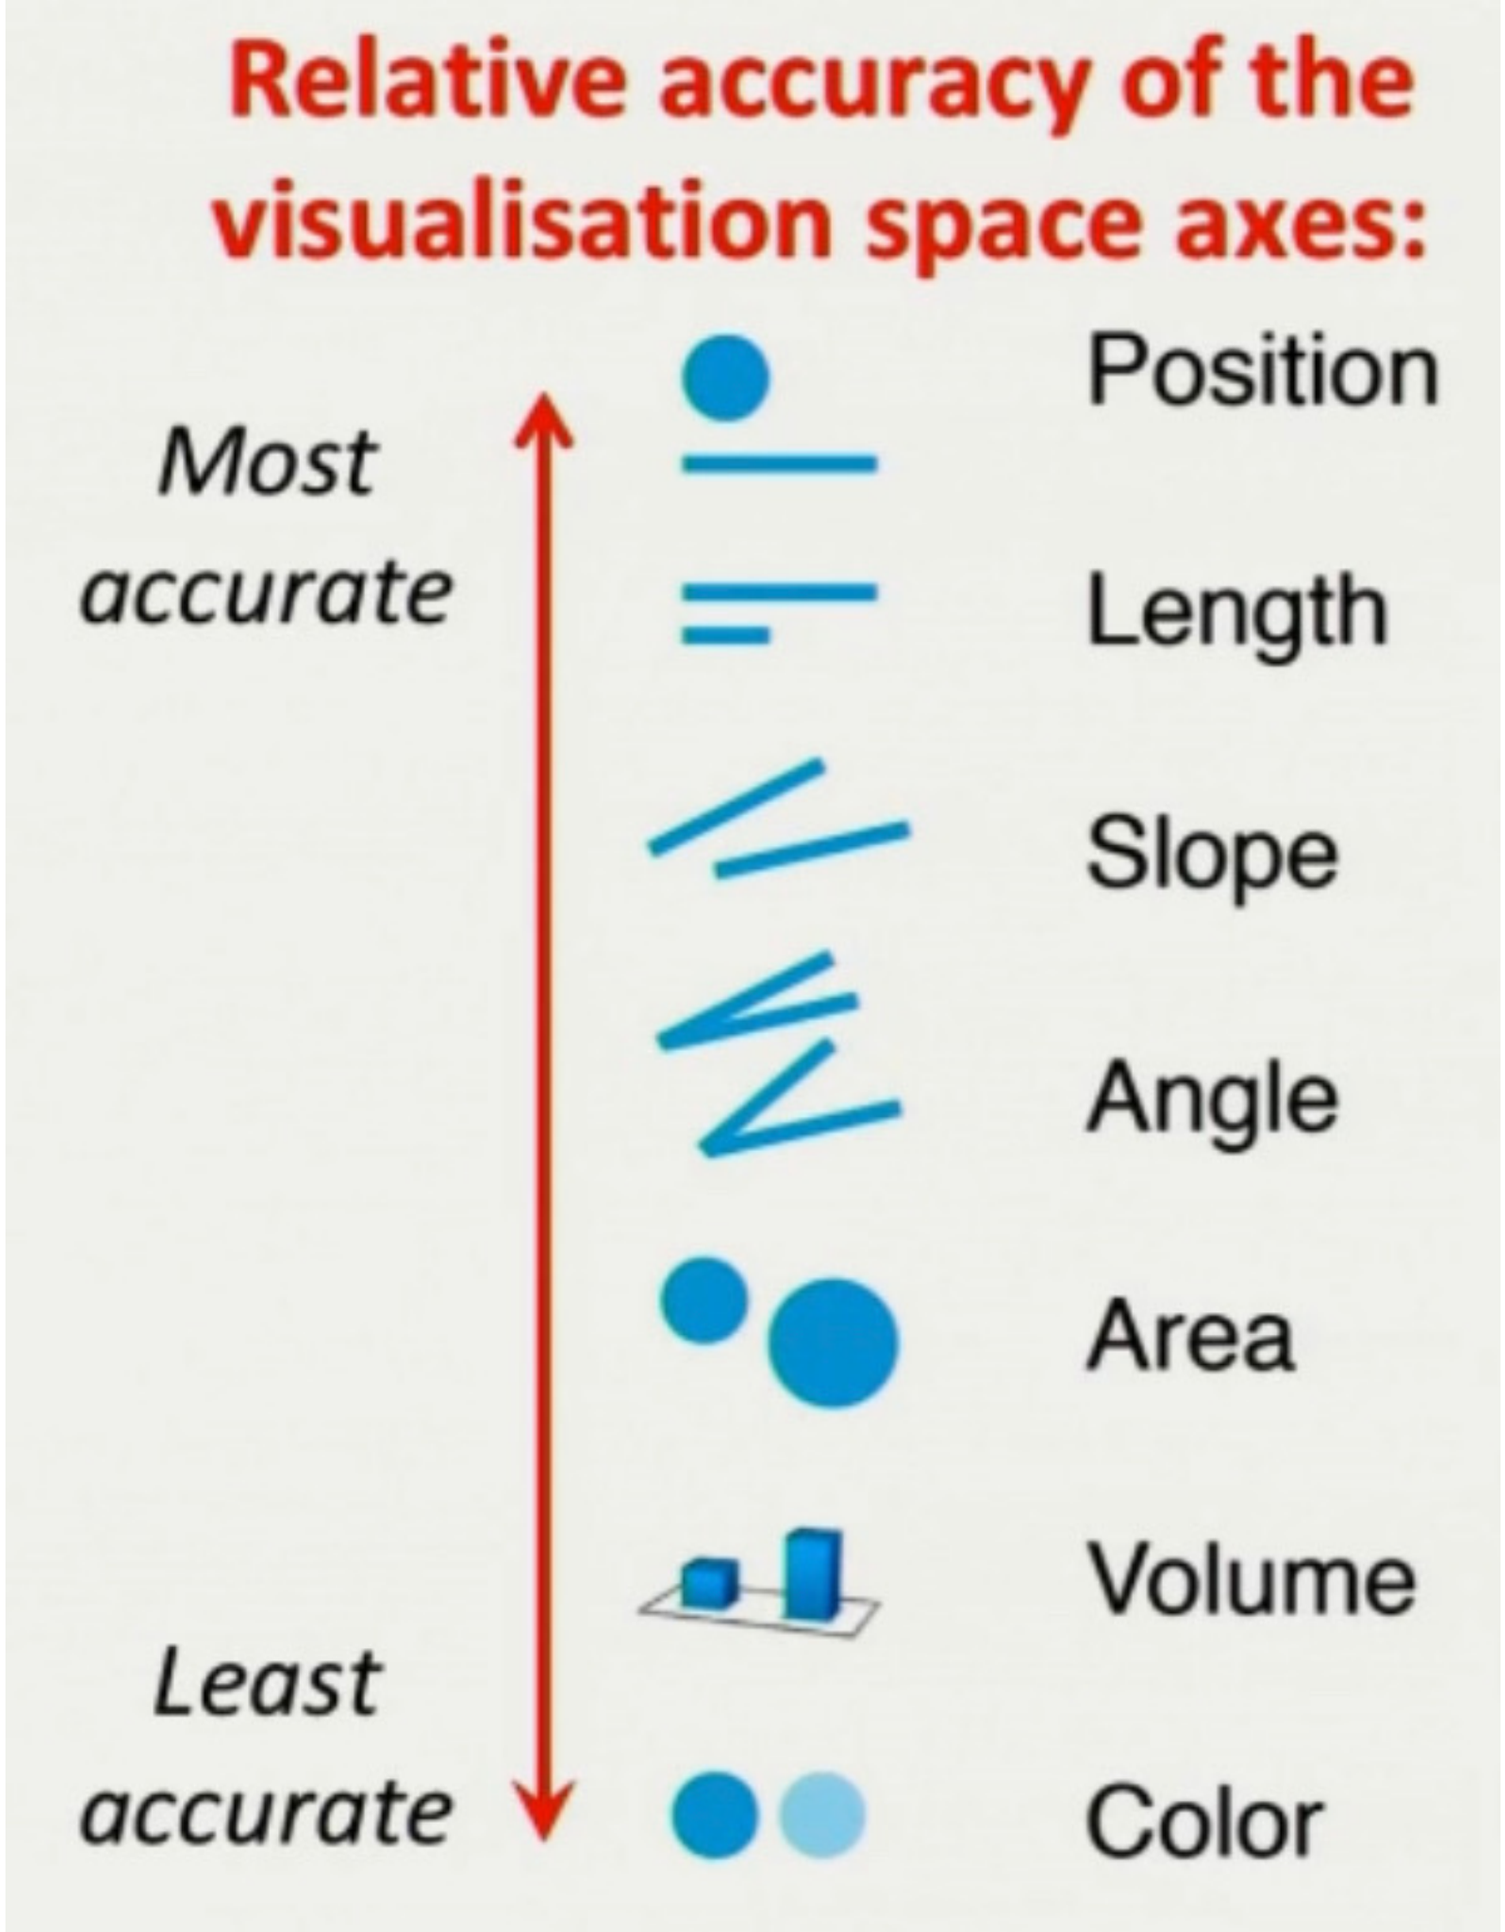
Relative accuracy of the visualisation space axes:
Most accurate
Position
Length
Slope
Angle
Area
Volume
Color
Least accurate

Hands on Sessions

- Galaxy/Star classification (with PCA)
- Regression Photometric redshift with Random Forests
- “voodoo lecture” on Deep Learning with Convolutional Neural Networks on denoising/transient finding

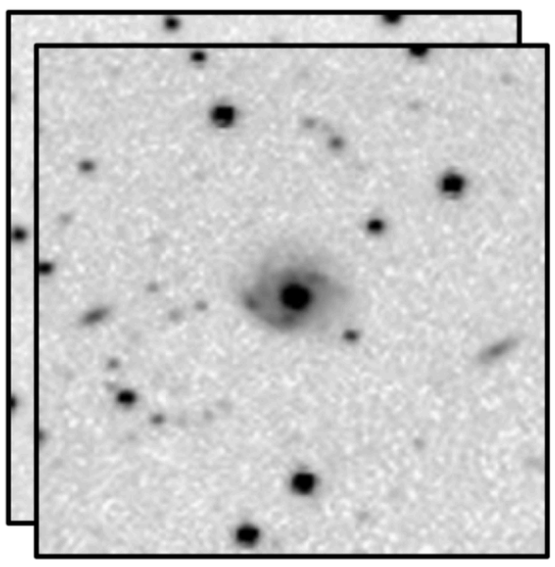

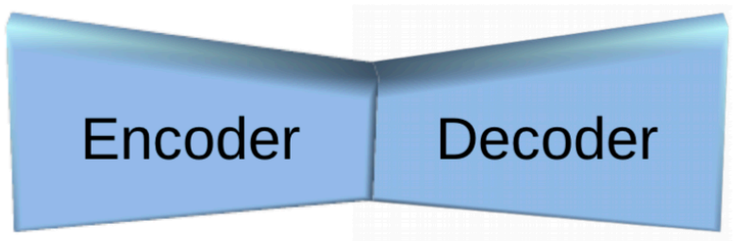

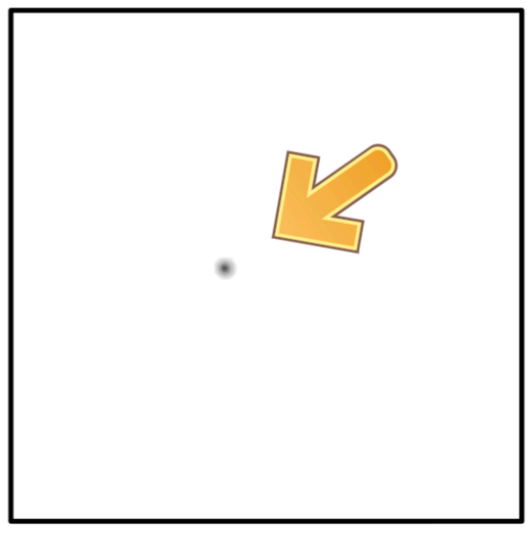

Name	Use	Logo
Pandas	Data Analysis	
Spark	Distributed Computing	
Scikit-learn	Machine Learning Toolbox	
Keras	Deep Learning	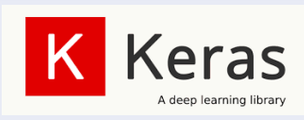
TensorFlow	Deep Learning	
Open-cv	Computer Vision	

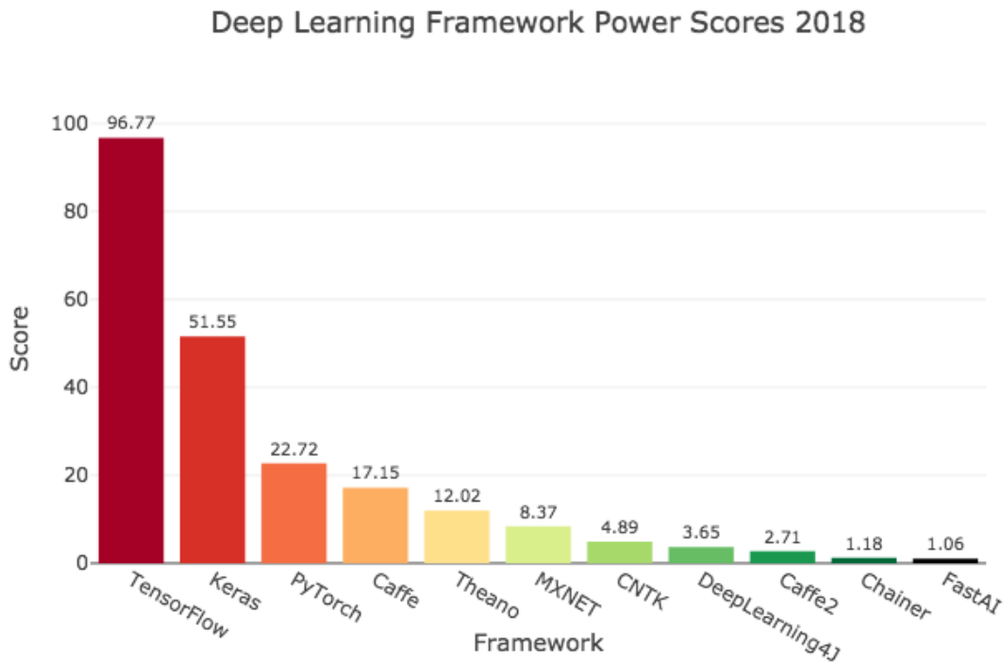
Deep Learning Framework Power Scores 2018
Score
Framework
TensorFlow
Keras
PyTorch
Caffe
Theano
MXNET
CNTK
DeepLearning4J
Caffe2
Chainer
FastAI
96.77
51.55
22.72
17.15
12.02
8.37
4.89
3.65
2.71
1.18
1.06


Amazon Books
Books
Framework
TensorFlow
Keras
PyTorch
Caffe
Theano
MXNET
CNTK
DeepLearning4J
Caffe2
FastAI
Chainer

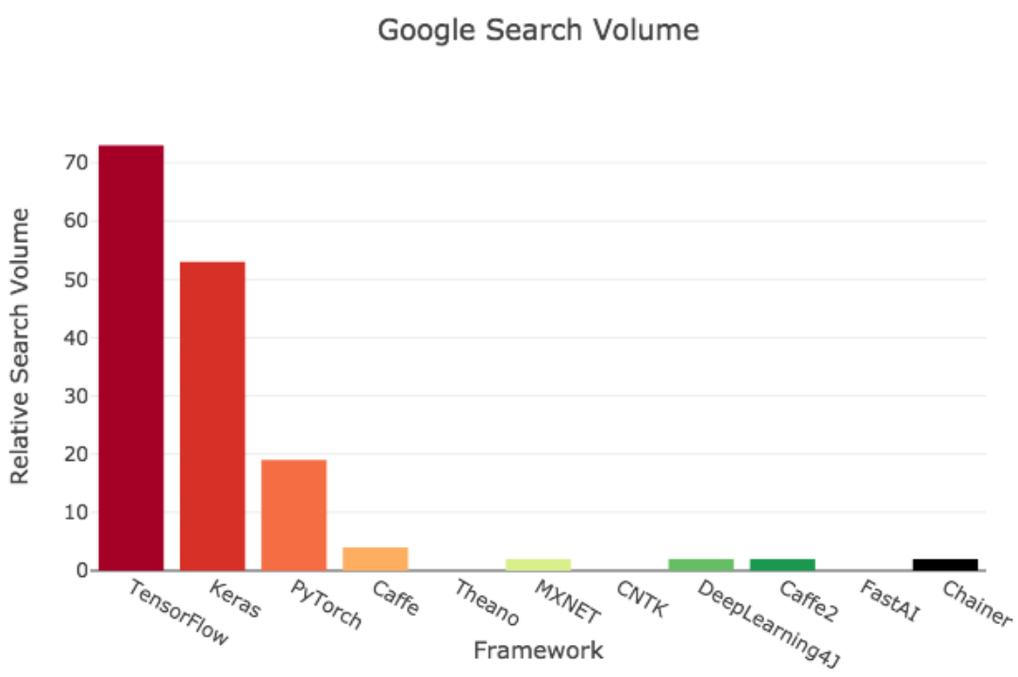
Google Search Volume
Relative Search Volume
Framework
TensorFlow
Keras
PyTorch
Caffe
Theano
MXNET
CNTK
DeepLearning4J
Caffe2
FastAI
Chainer

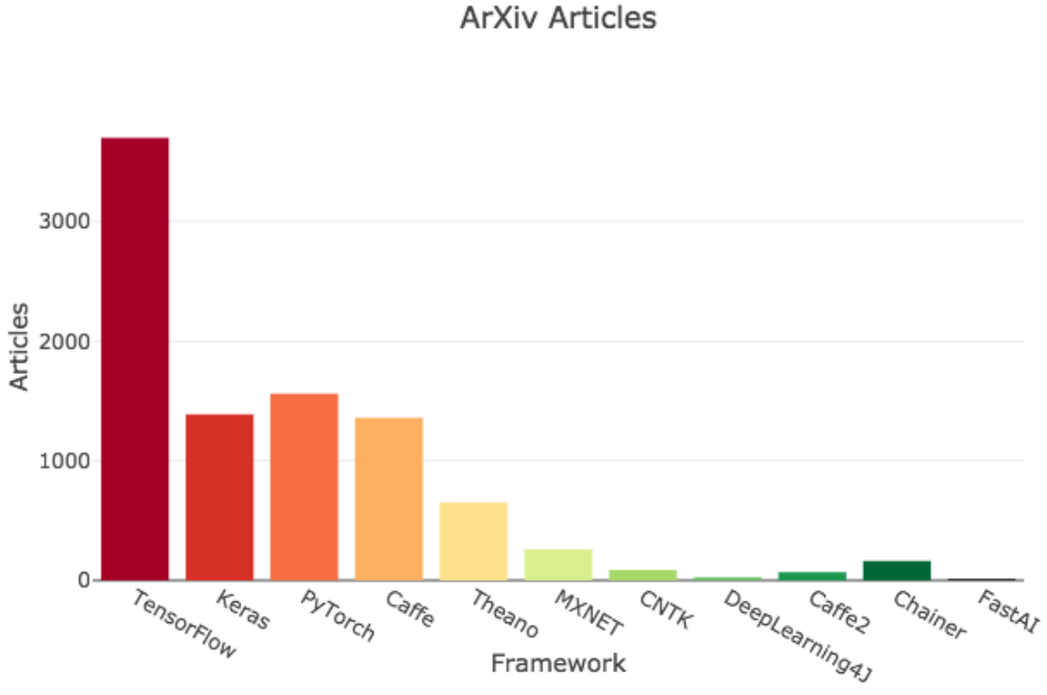
ArXiv Articles
Articles
Framework
TensorFlow
Keras
PyTorch
Caffe
Theano
MXNET
CNTK
DeepLearning4J
Caffe2
Chainer
FastAI

Software prerequisites

- Python 3
- Anaconda anaconda.org with Python 3.7
- You may create an environment
  `conda create –name=ml python=3.6`
- *scikit-learn*, *pandas*, *scikit-image*, *seaborn*, *tensorflow*, *keras* (may induce a downgrade to Python 3.6, but it's ok)
  `conda install xxx`
- Cheat sheets…

ML seminar

- Aka astro-ph
- an algorithm
- publications?
- 1x every 2 weeks or month?